

Testing the Deep Autoencoding Gaussian Mixture Mixture Model




Figure 2: An overview on Deep Autoencoding Gaussian Mixture Model

0.1 First Step: Dimension Reduction

The network takes as input a sample x and operates as shown below:

$$\mathbf{z}_c = h(\mathbf{x}; \theta_e)$$

$$\mathbf{x}' = g(\mathbf{z}_c; \theta_d)$$

$$\mathbf{z}_r = f(\mathbf{x}, \mathbf{x}')$$

$$\mathbf{z} = [\mathbf{z}_c, \mathbf{z}_r]$$

Here z_c is the low-dimensional embedding and x' is the reconstructed form of x produced by the encoder $h(\cdot)$ and decoder $g(\cdot)$. The variables θ_e and θ_d are parameters of the autoencoder. The function $f(\cdot)$ computes the reconstruction feature z_r . Observe that the value that will be passed to the estimation network is a matrix that is a combination of z_c and

0.2 Second Step: Density Estimation

Next, the equivalence between log-linear models and the posterior form of a Gaussian is exploited to perform density estimation via a simple feed-forward network which is referred to as the estimation network. The estimation network acts as a softmax classifier, providing a K-dimensional vector representing the membership probabilities:

$$\hat{\gamma} = \text{softmax}(\text{MLN}(\mathbf{z}; \theta_{\mathbf{m}})) \quad (1)$$

where θ_m is the parameter of the MLN.

Next, the parameters are adjusted following the same calculations as the typical EM algorithm.

$$\begin{aligned} \hat{\phi}_k &= \sum_{i=1}^N \frac{\hat{\gamma}_{ik}}{N} \\ \hat{\mu}_k &= \frac{\sum_{i=1}^N \hat{\gamma}_{ik} \mathbf{z}_i}{\sum_{i=1}^N \hat{\gamma}_{ik}} \\ \hat{\Sigma}_k &= \frac{\sum_{i=1}^N \hat{\gamma}_{ik} (\mathbf{z}_i - \hat{\mu}_k) (\mathbf{z}_i - \hat{\mu}_k)^T}{\sum_{i=1}^N \hat{\gamma}_{ik}} \\ E(\mathbf{z}) &= -\log \left(\sum_{k=1}^K \hat{\phi}_k \frac{\exp \left(-\frac{1}{2} (\mathbf{z} - \hat{\mu}_k)^T \hat{\Sigma}_k^{-1} (\mathbf{z} - \hat{\mu}_k) \right)}{\sqrt{|2\pi \hat{\Sigma}_k|}} \right) \end{aligned}$$


0.3 Objective Function

Finally we train the entire neural net with this somewhat mysterious objective function.

$$J(\theta_e, \theta_d, \theta_m) = \frac{1}{N} \sum_{i=1}^N L(\mathbf{x}_i, \mathbf{x}'_i) + \frac{\lambda_1}{N} \sum_{i=1}^N E(\mathbf{z}_i) + \lambda_2 P(\hat{\Sigma})$$

0.4 Explaining Quick Experiment

To get some intuition on how well the architecture is able to approximate the original distribution of the input I began by first coming up with my distributions with constant μ_k and Σ_k . An example is pictured below, depicting many points sampled from three Gaussians.



This example is somewhat unusual, and almost defeats the purpose of modeling the scenario as a GMM and was mostly done to see how well the model performs in certain edge-cases.

After this I feed the inputs into the neural-net. This required some tuning of the hyper parameters λ_1 and λ_2 . After training the model, I hope to find a suitable benchmark for comparing how well the model was able to reproduce the original distribution. For this I was considering simply implementing some standard clustering algorithms and comparing results.